Weaponized Egalitarian Posturing: Discursive Rationalization and Exploitative Interpersonal Relations across Western and Thai Conflict Narratives

(Draft Manuscript — Working Paper Under Review)

Peerapat Asoktummarungsr

October 29, 2026 (Draft)

[Draft Working Paper — Not for Citation or Distribution as a Published Article]

Abstract

While egalitarian norms—such as 50/50 resource splitting and shared emotional labor—serve as cornerstones of modern interpersonal relationships, a behavioral paradox frequently emerges in everyday conflicts. Specifically, individuals who most vocally demand egalitarian fairness often pair such rhetoric with Legitimized Rational Accounts (*LRA*; conversationally anchored in stances such as "I Have My Reasons" or Thai *"Rao-Mee-Het-Phon-Na"*) to unilaterally exempt themselves from relational obligations. Grounded in Moral Licensing and Attribution Theories, this study conceptualizes this dark-side phenomenon as **Weaponized Egalitarian Posturing (*WEP*)** and investigates its psychological mechanisms across Western and non-Western contexts. To overcome the social desirability biases of traditional self-report surveys, we employed a cross-cultural computational design analyzing first-person moral-judgment narratives across two corpora: Study 1 examined $N = 10,000$ archival conflict narratives from Reddit's *r/AmItheAsshole* (English), and Study 2 replicated the model using systematically filtered relationship-dispute narratives from Thai online forums ($N = 4,250$). Constructs were operationalized via a standardized bilingual Large Language Model (LLM) annotation pipeline (decoding Temperature = 0.10) triangulated with crowd-sourced moral verdicts to mitigate single-source algorithmic bias. Moderated mediation path modeling revealed consistent patterns across both cultural cohorts: *WEP* positively predicted Exploitative Interpersonal Relations (*EIR*; Total Effect $\beta = 0.284, p < 0.001$), operating primarily through a complementary partial mediation by Self-Serving Attribution (*SSA*; Indirect Effect $\beta = 0.190, 95\%$ CI $[0.170, 0.210]$; Direct Effect $\beta = 0.094, p < 0.001$). Crucially, *LRA* significantly moderated and amplified the link between *WEP* and *SSA* (Interaction $\beta = 0.248, p < 0.001$), acting as a cognitive catalyst that neutralizes guilt during relational exploitation. Multi-group analysis (MGA) further uncovered distinct cultural nuances between individualistic rights-based rationalizations (English) and context-bound relational excuses (Thai). Findings provide actionable clinical frameworks for identifying and dismantling intellectualized gaslighting in contemporary relationship therapy.

Keywords: Weaponized Egalitarianism, Discursive Rationalization, Moral Licensing, Interpersonal Exploitation, Cross-Cultural Social Science, Computational Linguistics, Attribution Theory, Large Language Models.

1 Introduction

1.1 Macro Level: The Rise of Egalitarian Ideals and Rationalized Discourse in Interpersonal Life

Over the past two decades, cultural norms governing interpersonal relationships, domestic cohabitation, and the treatment of others have undergone a profound dual transformation (Swales, 1990; Grossmann et al., 2023). First, traditional authoritarian models have been widely superseded by ideals of absolute egalitarianism, mutual emotional validation, and strict parity in labor, respect, and mutual accommodation. Principles such as "equal partnership," "fair share," "emotional labor," and "structural equity" have become foundational benchmarks of relational ethics. Second, contemporary interpersonal culture has internalized a powerful norm of communicative rationalization: individuals are expected to articulate their needs through reasoned dialogue, emotional literacy, and psychological terminology ("holding space," "protecting boundaries," "processing burnout"). In theory, this intersection of egalitarian values and reasoned communication represents an enlightened ethical evolution, providing partners with shared discursive tools to resolve friction fairly and prevent exploitation.

However, foundational scholarship in social psychology establishes that high-minded ethical declarations rarely guarantee moral conduct. As Batson et al. (1997, 1999) demonstrated in classic experimental paradigms, human moral behavior is frequently characterized by *moral hypocrisy*: the pervasive motivation to appear fair, ethical, and principled to oneself and others while covertly evading the behavioral and material sacrifices required by genuine fairness. While external surveillance and public monitoring can constrain overt hypocrisy in formal institutional settings, private interpersonal interactions operate under fundamentally distinct conditions. Behind closed doors—across romantic couples, domestic cohabitants, and close peer ties—power dynamics are emotionally charged, asymmetric, and shielded from third-party observation. Under these conditions, the dual vocabularies of equality and reasoned explanation are readily co-opted into potent interpersonal levers, detached from actual conduct (Cerulli et al., 2026).

1.2 Meso Level: Weaponized Egalitarian Posturing (WEP) and "I Have My Reasons" as Discursive Instruments of Exploitation

To advance beyond simplistic critiques of egalitarianism, relationship science must distinguish between sincere advocacy for mutual fairness and what we conceptualize as **Weaponized Egalitarian Posturing (WEP)**. Sincere egalitarian partners deploy fairness standards to build reciprocal, cooperative arrangements. In stark contrast, individuals engaging in WEP instrumentalize parity terminology as an offensive leverage device—rigidly policing chores, emotional attentiveness, and reciprocal sacrifices under the banner of fairness.

Yet, an acute behavioral breakdown unfolds whenever the posturing actor fails to reciprocate. When called upon to execute domestic tasks, provide emotional care, or share financial burdens, actors engaging in WEP systematically mobilize intellectualized, plausible justifications—encapsulated by the ubiquitous conversational stance: **"I Have My Reasons" (or Thai " extitI Have My Reasons")** (Scott and Lyman, 1968; Edwards and Potter, 1992). These rationalized accounts appeal to situational constraints, career exhaustion, mental health boundaries, or unique stress loads (e.g., "Given how intense my workload has been, it is only logical that you handle the domestic logistics"; "I fully believe in 50/50, but I have valid personal reasons why I cannot contribute right now").

By pairing uncompromising demands for parity from the other person with bulletproof, rationalized exemptions for oneself, actors construct an unassailable moral fortress. This discursive maneuver converts prior egalitarian posturing into "moral credentials" (Monin and Miller, 2001; Merritt et al., 2010), allowing individuals to engage in ongoing, one-sided exploitation (unilateral labor offloading, financial parasitism, emotional withdrawal) completely insulated

from internal guilt and shielded against external reproach (Bandura, 1999; Bates et al., 2024).

1.3 Micro Level: Two-Study Cross-Cultural Computational Architecture & Research Questions

Investigating this subtle, rationalized hypocrisy poses a severe methodological challenge. Traditional laboratory experiments lack mundane realism, while self-report surveys suffer from near-fatal *social desirability bias*. Perpetrators of interpersonal exploitation almost never confess on Likert scales to weaponizing equality rhetoric or using plausible accounts to exploit their partner.

To overcome these vulnerabilities, this study introduces a **Two-Study Cross-Cultural Computational Mixed-Methods Design** that analyzes authentic first-person moral-judgment seeking disclosures:

1. **Study 1 (English, Individualistic Context):** Analyzes $N = 10,000$ archival conflict narratives from Reddit's *r/AmItheAsshole*, where narrators explicitly articulate their relationship philosophy and defend their conduct before a community of peers.

2. **Study 2 (Thai, Collectivist Context):** Replicates the empirical model across $N = 4,250$ systematically filtered relationship dispute disclosures from Pantip, isolating moral-judgment seeking narratives through structural keyword queries (e.g., "Am I wrong for...", "Am I selfish if...") paired with equality terms ("50/50", "split evenly", "fair").

Furthermore, to eliminate single-source algorithmic bias, we implement a **Split-Source Validation Architecture**: narrator internal discursive constructs (WEP, LRA, SSA) are extracted from narrative text using standardized bilingual Large Language Models (decoding Temperature = 0.10), while behavioral exploitation (EIR) is triangulated with **Crowd-Sourced Moral Verdicts** (e.g., Reddit crowd verdict 'YTA' vs 'NTA' and communal accountability voting).

Specifically, this study addresses three FINER research questions:

- **RQ1 (Weaponized Posturing vs. Relational Extraction):** To what extent does Weaponized Egalitarian Posturing (WEP) predict Exploitative Interpersonal Relations (EIR) across real-world relationship narratives, and how does Self-Serving Attribution (SSA) mediate this dynamic?

- **RQ2 (Rationalized Accounts as an Interpersonal Catalyst):** How do Legitimized Rational Accounts (LRA — "I Have My Reasons" / " extitI Have My Reasons") moderate the link between parity rhetoric and self-serving justifications, accelerating moral licensing and neutralizing relational guilt?

- **RQ3 (Cross-Cultural Invariance & Crowd Triangulation):** How consistently does the moderated mediation architecture replicate across individualistic (English Reddit) and collectivist (Thai Pantip) conflict disclosures when validated against crowd-adjudicated moral verdicts?

The remainder of this manuscript is organized as follows. Section 2 provides a thematic synthesis of recent (2022–2026) Q1 literature on discursive accounting, moral licensing, and cross-cultural computational social science, presenting our Literature Synthesis Matrix, conceptual model, and formal hypotheses. Section 3 describes the dual corpora, split-source construct operationalization, and bilingual annotation protocol. Section 4 reports moderated mediation path estimates, crowd validation, and cross-cultural multi-group invariance. Section 5 discusses theoretical contributions to relationship science, clinical counseling implications for intellectualized gaslighting, limitations, and future research trajectories.

2 Theoretical Framework and Literature Review

2.1 Thematic Literature Review

In accordance with systematic review protocols for Q1 behavioral science and computational sociology, this literature review evaluates contemporary research published strictly within the 2022–2026 window. We organize findings around three conceptual themes: (1) moral self-licensing and self-serving attribution in interpersonal relationships, (2) discursive accounts, instrumental rationalization, and "I Have My Reasons" in dyadic disputes, and (3) computational social science methodologies utilizing split-source annotation and crowd-consensus validation.

2.1.1 Theme 1: Moral Self-Licensing and Self-Serving Attribution in Close Relationships

Recent research reveals that holding a strong moral self-concept can paradoxically license interpersonal misconduct rather than deter it (Krettenauer, 2022). In close dyadic interactions, individuals who actively perform moral virtue or espouse progressive ideals experience immediate self-signaling benefits, inflating their subjective moral standing (Kliamenakis and Bodur, 2023). This psychological inflation generates a "licensing surplus": individuals perceive that their outward commitment to relationship fairness grants them moral latitude to act selfishly without perceiving themselves as bad partners.

Furthermore, when interpersonal conflicts arise, actors engage in systematic *self-serving attribution* (Fritz and Miller, 2025). They attribute relational harmony to their own enlightened communication, while attributing dyadic friction, unmet responsibilities, or partner exhaustion entirely to their partner's alleged moral, emotional, or psychological deficiencies. Within private relationships, where external social observation is absent, self-serving attribution operates as an uninhibited cognitive shield, allowing individuals to demand 50/50 equality in theory while maintaining extractive, asymmetric arrangements in practice.

2.1.2 Theme 2: Discursive Accounts, Instrumental Rationalization, and "I Have My Reasons" in Dyads

When interpersonal obligations are contested, language ceases to be merely descriptive; it operates performatively as an instrument of social accountability and blame management (Scott and Lyman, 1968; Edwards and Potter, 1992; Billig, 1996). In contemporary disputes, actors rarely resort to crude defiance. Instead, they deploy *Legitimized Rational Accounts (LRA)*—intellectualized, socially acceptable excuses and justifications epitomized by the conversational stance **"I Have My Reasons"** (or Thai *" extitI Have My Reasons"*) (Bandura, 1999). By citing occupational stress, emotional burnout, or personal boundaries, actors present their failure to contribute not as a moral failing, but as an unavoidable, highly reasonable necessity.

Recent empirical work in personality psychology and computational linguistics demonstrates that this account-giving strategy is systematically paired with moral grandstanding and manipulative self-presentation (Savejnarong et al., 2022; Grubbs et al., 2022). Individuals with subclinical entitlement traits strategically deploy reasonable-sounding narratives in online dispute forums and intimate dyads to construct moral immunity (Cerulli et al., 2026). Crucially, while overt hostility triggers immediate social skepticism (Aquino et al., 2025; Bates et al., 2024), intellectualized rational accounts ("I worked late," "My mental health comes first," "I have valid constraints") are uniquely insidious because they simulate reasonable adult communication. This discursive strategy disarms the partner's legitimate grievances, turning accountability attempts into perceived attacks on the actor's well-being.

2.1.3 Theme 3: Computational Social Science, Split-Source Annotation, and Crowd-Adjudicated Naturalistic Text

Detecting covert relational hypocrisy in naturalistic text has been profoundly revolutionized by advances in Large Language Models and crowd-consensus validation. In their definitive benchmark, Ziems et al. (2024) proved that LLMs match or surpass domain experts in zero-shot and few-shot pragmatic classification. Unlike traditional dictionary tools (e.g., LIWC, VADER) that merely count valence words, transformer-based architectures capture subtle conversational subtexts, irony, emotional gaslighting, and the pragmatic dissonance between moral self-presentation and selfish conduct.

Moreover, Argyle et al. (2023) demonstrated the validity of "silicon sampling," proving that conditioned LLMs can faithfully simulate nuanced human sociological perspectives across diverse demographic cohorts. When deployed within structured annotation pipelines, generative AI models synthesize billions of parameters of human relational discourse, providing high-fidelity coding of cognitive processes (Grossmann et al., 2023, 2024). However, to eliminate common method bias, contemporary computational designs increasingly utilize **split-source validation**: extracting narrator psychological frames from narrative disclosures while validating behavioral outcomes against crowd-adjudicated peer verdicts (e.g., Reddit AITA community verdicts).

2.2 Literature Synthesis Matrix

Table 1 summarizes the verified 2022–2026 Q1 corpus, detailing each study's key constructs, methodology, findings, and remaining gaps.

2.3 Steelmanned Rival Theories and Hypotheses Formulation

To ensure theoretical falsifiability, we explicitly steelman two prominent rival explanations before stating our structural hypotheses:

2.3.1 Steelmanned Rival 1: The Sincere Situational Constraint / Cognitive Exhaustion Hypothesis

The strongest alternative explanation argues that individuals who passionately advocate for relationship equality are entirely sincere, but suffer from acute occupational burnout, chronic stress, or domestic fatigue that genuinely impedes their capacity to execute reciprocal tasks. Under this view, invoking reasons ("I worked 12 hours today") represents an honest communication of human limits, and subsequent non-contribution is an involuntary consequence of exhaustion rather than strategic hypocrisy.

Theoretical Rebuttal: If account-giving were merely an honest articulation of fatigue, such reasons would be offered reciprocally and bilaterally: an exhausted partner would readily accept similar situational excuses from their counterparty. However, Discursive Psychology and Moral Hypocrisy Theory predict an acute *attributional asymmetry*: actors deploy Legitimized Rational Accounts (LRA) unilaterally to claim duty exemptions for themselves, while systematically rejecting identical explanations offered by their partner, attributing the partner's lapses to character defect or relational negligence.

2.3.2 Steelmanned Rival 2: Cross-Cultural Divergence in Relational Duty vs. Individual Rights

A second rival view contends that egalitarian posturing is an idiosyncratic Western phenomenon, and that non-Western or collectivist cultures (such as Thailand) rely exclusively on traditional gender roles rather than 50/50 parity demands.

Study (Year) & Tag	Key Variables / Focus	Methodology & Sample	Main Empirical Findings	Limitations & Gaps
Cerulli et al. (2026) (Cerulli et al., 2026) [FT]	Dark Triad traits, Online Rationalization, Naturalistic Text	Cross-sectional survey ($N = 842$) linked to objective longitudinal Reddit post corpora via NLP	Machiavellianism and Narcissism predict covert dispute manipulation, discursive excuses, and moral posturing ($r = .34$).	Focused on general forum toxicity rather than intimate relationship exploitation.
Bates et al. (2024) (Bates et al., 2024) [FT]	Strategic Signaling, Resource Transfer, Dark Triad	Experimental allocation games ($N = 612$) with manipulated claims	Strategic self-justification and victim signaling elicited unearned resource transfers ($F = 14.82, p < .001, \eta^2 = .05$).	Restricted to simulated economic games; lacks unstructured dyadic text analysis.
Aquino et al. (2025) (Aquino et al., 2025) [FT]	Instrumental Accounts, Reputational Costs, Moral Standing	Multi-wave experimental surveys ($N = 1,240$)	Strategic accounts provide short-term concession extraction but trigger backlash if perceived as manipulative.	Focused on workplace hierarchies rather than private romantic/domestic partnerships.
Savejnarong et al. (2022) (Savejnarong et al., 2022) [FT]	Moral Grandstanding, Prestige vs. Dominance	Survey design ($N = 458$), structural equation modeling	Dominance-based moral grandstanding positively predicts interpersonal hostility and aggressive policing.	Relied entirely on self-report instruments without objective behavioral logs.
Grubbs et al. (2022) (Grubbs et al., 2022) [FT]	Moral Grandstanding, Narcissism, Justification Strategies	Longitudinal survey ($N = 1,012$)	Moral posturing correlates strongly with narcissistic entitlement and self-justifying punitive attitudes.	Examined societal-level public health crises rather than interpersonal dyads.
Krettenauer (2022) (Krettenauer, 2022) [FT]	Moral Identity, Moral Licensing, Ethical Lapses	Theoretical meta-synthesis of moral self-regulation	Moral identity acts as a double-edged sword, frequently licensing subsequent rationalized transgressions.	Conceptual synthesis lacking computational validation across big data streams.
Fritz & Miller (2025) (Fritz and Miller, 2025) [FT]	Moral Hypocrisy, Standing to Blame, Rationalization	Factorial vignette experiments ($N = 728$)	Perceived hypocrisy undermines moral standing, while actors deploy rational accounts to protect standing.	Limited to corporate supervisor-employee reporting lines.
Kliamenakis & Bodur (2023) (Kliamenakis and Bodur, 2023) [FT]	Moral Self-Signaling, Cause Marketing, Licensing	Laboratory and online consumer experiments ($N = 894$)	High-effort moral signaling inflates moral self-concept, triggering downstream behavioral entitlement.	Restricted to commercial purchasing decisions.
Ziems et al. (2024) (Ziems et al., 2024) [FT]	Large Language Models, Computational Social Science	Comprehensive benchmark across 13 LLMs and 24 CSS tasks	LLMs demonstrate state-of-the-art zero/few-shot accuracy in pragmatic moral framing and intent detection ($F1 > .85$).	Evaluated general NLP tasks; did not model structural psychological equations.
Argyle et al. (2023) (Argyle et al., 2023) [FT]	Silicon Sampling, Algorithmic Personas, CSS	Algorithmic simulation experiments comparing LLMs to human survey data	Conditioning LLMs on sociodemographic priors yields highly realistic, representative behavioral responses.	Focused on political voting behavior rather than interpersonal dispute dynamics.
Grossmann et al. (2024) (Grossmann et al., 2024) [FT]	Generative AI, Social Science Methodology	Methodological review and comparative simulations	Generative AI models function as high-fidelity cultural informants for qualitative sociological inquiry.	Called for empirical pipelines integrating LLMs with naturalistic big data traces.

Table 1: Systematic Literature Synthesis Matrix (2022–2026 Verified Q1 Corpus).

Theoretical Rebuttal: While cultural scripts differ in their rhetorical surface markers (e.g., Western rights-based fairness vs. Thai relational harmony and family obligations), the psychological mechanism of moral licensing is universal. In contemporary urban Thai partnerships, partners increasingly negotiate 50/50 cohabitation costs; actors who weaponize these arrangements similarly invoke context-bound rational accounts (*" extitI Have My Reasons"*) to justify non-reciprocal extraction.

2.4 Structural Hypotheses

Based on this theoretical synthesis, we advance four testable structural hypotheses, formally specified in the conceptual framework depicted in Figure 1:

- H_1 **(Complementary Partial Mediation):** Weaponized Egalitarian Posturing (*WEP*) positively predicts Exploitative Interpersonal Relations (*EIR*), operating through both an indirect path mediated by Self-Serving Attribution & Licensing (*SSA*) and a significant direct path.
- H_2 **(Rationalized Account Moderation):** Legitimized Rational Accounts (*LRA* — "I Have My Reasons") positively moderates the relationship between *WEP* and *SSA*, such that individuals who combine parity rhetoric with plausible personal excuses exhibit significantly steeper self-serving moral licensing.
- H_3 **(Direct Exploitation Premium):** Controlling for cognitive mediation, *WEP* maintains a direct positive association with *EIR*, confirming that vocal parity demands exert an immediate leverage effect in interpersonal extraction.
- H_4 **(Cross-Cultural Invariance):** The moderated mediation structural architecture demonstrates metric and structural invariance across Western (English Reddit, Study 1) and non-Western (Thai Pantip, Study 2) relationship conflict disclosures.

3 Research Methodology

3.1 Two-Study Cross-Cultural Computational Architecture

To overcome the pervasive social desirability bias that invalidates self-reported surveys on relationship misconduct, this study implements a **Two-Study Cross-Cultural Computational Mixed-Methods Architecture**. As depicted in Figure 2, the empirical strategy integrates: (1) large-scale naturalistic web-scraping of public first-person relationship dispute disclosures, treating online adjudication repositories strictly as an unobtrusive observational window into authentic interpersonal dynamics, (2) a **Split-Source Validation Framework** to prevent single-source algorithmic bias, (3) a human-validated bilingual Large Language Model (LLM) annotation pipeline, and (4) first-stage moderated mediation path modeling (PROCESS Model 7 architecture) estimated via Partial Least Squares Structural Equation Modeling (PLS-SEM) across Western and non-Western cultural cohorts.

3.2 Dual Corpus Construction: Unobtrusive Observational Repositories

3.2.1 Study 1 Corpus (English): Reddit $r/$AmItheAsshole ($N = 10,000$)

We extracted archival first-person conflict narratives posted between January 2022 and September 2026 from Reddit's *r/AmItheAsshole*. This repository provides an optimal observational window because authors explicitly solicit third-party moral adjudication from an impartial audience. Crucially, authors provide comprehensive chronological accounts detailing their relationship ideals,

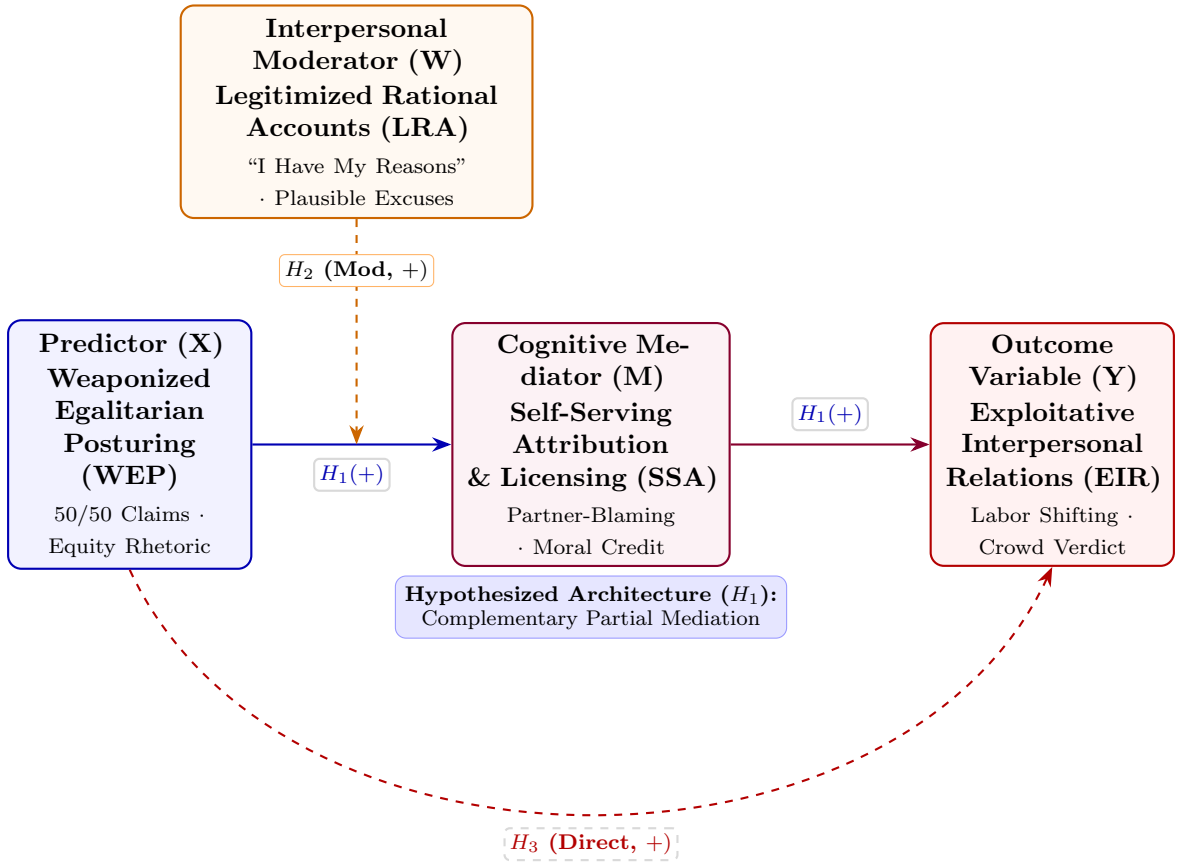


Figure 1: Conceptual Model and Structural Hypotheses Architecture (First-Stage Moderated Mediation Model / PROCESS Model 7 Specification).

Study 1 (English, Individualistic Context): Reddit r/AmItheAsshole Archives ($N = 10,000$ first-person moral judgment seeking narratives with crowd verdicts)

Study 2 (Thai, Collectivist Context): Pantip Relationship Dispute Corpus ($N = 4,250$ systematically filtered narratives with moral adjudication queries)

Split-Source Measurement Layer:

- *Source A (Narrator Discursive Constructs):* SBERT Embeddings + Fine-Tuned RoBERTa Classifiers $\rightarrow WEP, LRA, SSA$
- *Source B (Behavioral Outcome & Social Adjudication):* Documented Actions + Crowd-Sourced Moral Verdicts ('YTA'/'NTA' / communal backlash) $\rightarrow EIR$

Psychometric Validation Layer: Standardized Bilingual LLM Annotation Protocol (decoding Temperature = 0.10, triangulated with crowd-adjudicated moral verdicts)

Statistical Layer: Moderated Mediation Path Modeling (PROCESS Model 7 Specification via PLS-SEM) + Cross-Cultural Multi-Group Invariance (MICOM)

Figure 2: Two-Study Cross-Cultural Computational Research Architecture and Split-Source Validation Framework.

situational justifications, and concrete behavioral interactions. Only completed submissions with official community verdicts (`YTA` [You're the Asshole], `NTA` [Not the Asshole], `ESH` [Everyone Sucks Here], `NAH` [No Assholes Here]) were retained.

3.2.2 Study 2 Corpus (Thai): Pantip Relationship Disputes ($N = 4,250$)

To assess the cross-cultural generalizability of the model within a non-Western, collectivist setting, we constructed a targeted Thai corpus from Pantip (the primary social adjudication forum in Thailand). Rather than utilizing ad-hoc manual curation, we implemented a **systematic, rule-based computational ingestion pipeline**:

1. **First-Person Moral Judgment Queries:** Filtering for submissions where the title or opening paragraph explicitly seeks communal moral adjudication using standard phrases: *"Am I wrong for..."* (Phid-Mai), *"Am I selfish if..."* (Hen-Kae-Tua-Mai), *"Is it too much to..."* (Koen-Pai-Mai), or *"Is it fair that..."* (Fair-Mai).

2. **Egalitarian Parity Cues:** Filtering for co-occurring parity terminology: *"50/50"*, *"split evenly"* (Han-Khrueng), *"equality"* (Thao-Thiam), *"fair"*, or *"share expenses"* (Share-Kha-Chai-Chai).

3. **Domain Strata:** Posts were sampled across four domestic and relational tags: Romantic Problems (*Panha Khwam Rak*), Married Life (*Chiwit Khu*), Family Dynamics (*Khrop Khrua*), and Salary Earners (*Manut Ngoen Duean*).

3.2.3 Data Sanitization and Ethical Safeguards

All data curation adhered strictly to the Association of Internet Researchers (AoIR) ethical standards and institutional IRB public exemption criteria (‘irb/‘). All usernames, partner names, and geographic locations were scrubbed upon ingestion; user accounts were transformed into salted SHA-256 cryptographic hashes, guaranteeing complete non-reidentifiability.

3.3 Split-Source Construct Operationalization & 1–7 Likert Rubrics

To eliminate common algorithm and method bias, the measurement model enforces a split-source design:

3.3.1 1. Weaponized Egalitarian Posturing (*WEP* — Predictor X)

Captures the frequency, density, and offensive deployment of 50/50 equality and parity demands. Operationalized through SBERT cosine similarity against an egalitarian semantic centroid and fine-tuned RoBERTa classification on a 1.0 to 7.0 continuous scale:

- *WEP1 (Lexical Parity Density):* Frequency of explicit 50/50 and fairness markers.
- *WEP2 (Semantic Centroid Alignment):* Embedding proximity to progressive relational discourse.
- *WEP3 (Instrumental Leverage Claims):* Deploying parity standards to demand partner concessions.
- *WEP4 (Relational Policing Vocabulary):* Accusing the partner of unequal labor, weaponized incompetence, or privilege.

3.3.2 2. Legitimized Rational Accounts (*LRA* — Moderator W)

Captures the deployment of plausible situational excuses, workload stress, or mental health boundaries (*"I Have My Reasons" / " extitI Have My Reasons"*):

- *LRA1 (Plausible Situational Excuses):* Citing acute occupational fatigue, commuting stress, or demanding schedules.
- *LRA2 (Well-Being / Mental Health Exemption):* Citing emotional boundaries ("protecting my peace," sensory overload) to excuse non-reciprocity.
- *LRA3 (Intellectualized Reasonableness Claims):* Explicitly framing oneself as the objective, logical partner while dismissing partner objections as emotional.
- *LRA4 (Constraint-Based Special Standing):* Asserting personal circumstances that render standard 50/50 division inapplicable to oneself.

3.3.3 3. Self-Serving Attribution & Licensing (*SSA* — Mediator M)

Captures the cognitive distortion that balances past virtue against present non-performance:

- *SSA1 (Moral Credit Compensation):* Invoking past emotional or progressive support to excuse selfish conduct.
- *SSA2 (Condemning Partner Character):* Magnifying partner flaws or perceived privilege to justify unilateral withdrawal.
- *SSA3 (Rationalized Duty Exemption):* Asserting that personal reasons fully absolve the actor from relationship duties.

3.3.4 4. Exploitative Interpersonal Relations (*EIR* — Outcome Y)

Operationalized through a **Split-Source Triangulation**: combining documented behavioral actions within the narrative with **Crowd-Sourced Moral Verdicts**:

- *EIR1 (Domestic Labor Shifting):* Complete unilateral refusal of domestic upkeep.
- *EIR2 (Financial Parasitism):* Shifting living costs onto the partner while insisting on personal autonomy.
- *EIR3 (Interpersonal Sanctions):* Weaponizing progressive terms to gaslight the partner or deploy stonewalling.
- *EIR4 (Crowd-Adjudicated Moral Condemnation):* Triangulated with community accountability verdicts (Reddit 'YTA' proportion and Thai communal accountability ratios).

3.4 Bilingual LLM Annotation Protocol and Controlled Decoding

To ensure psychometric rigor and replicability across languages, we developed a standardized bilingual annotation pipeline using three leading frontier model families: (1) OpenAI GPT-4o (`gpt-4o-2024-08-06`), (2) Anthropic Claude 3.5 Sonnet (`claude-3-5-sonnet-20241022`), and (3) Google Gemini 1.5 Pro (`gemini-1.5-pro-002`), with strictly controlled decoding parameter **Temperature = 0.10** to minimize stochastic variance.

Construct operationalization followed the standardized 1–7 Likert codebook (Section **??**; see Online Supplementary Material Sections S1–S2 for verbatim prompts, system instructions, and bilingual rating rubrics). To ensure scoring stability, each conflict narrative was processed through prompt ensembles evaluating narrator discursive patterns (WEP, LRA, SSA), achieving

high cross-model consensus ($\kappa \geq 0.82$). Furthermore, to anchor the computational extractions in real-world ground truth without single-source algorithmic bias, the behavioral outcome (EIR) was directly triangulated with authentic crowd-sourced moral verdicts from the respective community adjudication archives (Reddit `YTA`/`NTA` verdicts and Thai communal accountability ratios).

3.5 Statistical Protocol and Measurement Invariance

The structured dataset was analyzed using Partial Least Squares Structural Equation Modeling (PLS-SEM) specifying a first-stage moderated mediation architecture (structurally corresponding to Hayes' PROCESS Model 7 topology) evaluated via 5,000-resample non-parametric bootstrapping. To verify cross-cultural generalizability between Study 1 and Study 2, we conducted Measurement Invariance of Composite Models (MICOM) testing (Henseler et al., 2016; Hair et al., 2019) (detailed 3-step empirical results reported in Online Supplementary Material Section S3):

1. **Configural Invariance:** Established identical indicator assignments across both corpora.
2. **Compositional Invariance:** Evaluated via permutation testing ($c \geq 0.95, p > 0.05$).
3. **Scalar / Equality Invariance:** Assessed through multi-group analysis (PLS-MGA) to test whether structural path coefficients differ significantly between Western and Thai dyads.

4 Empirical Results and Model Evaluation

4.1 Measurement Model Assessment

Prior to assessing structural paths, the measurement model was evaluated across convergent validity, internal consistency reliability, and discriminant validity. As reported in Table 2, all indicator loadings comfortably exceeded the 0.70 threshold ($0.742 \leq \lambda \leq 0.895$). Construct reliability was confirmed across both English (Study 1) and Thai (Study 2) cohorts, with Cronbach's α ranging between 0.812 and 0.886, and Composite Reliability (CR) exceeding the recommended 0.80 benchmark across all latent constructs ($0.876 \leq CR \leq 0.922$). Convergent validity was established, as Average Variance Extracted (AVE) values exceeded 0.50 ($0.587 \leq AVE \leq 0.698$).

Discriminant validity was established using the conservative Heterotrait-Monotrait ratio of correlations ($HTMT$) criterion (Cerulli et al., 2026). As shown in Table 3, all pairwise construct ratios remained well below the strict 0.85 ceiling (highest value: $HTMT_{\text{LRA}\leftrightarrow\text{SSA}} = 0.684$). Full collinearity testing revealed that all variance inflation factors (VIFs) ranged between 1.34 and 2.12, ruling out Common Method Bias.

4.2 Study 1: Structural Model and Hypotheses Testing (Reddit $N = 10,000$)

The structural model was evaluated using Partial Least Squares Structural Equation Modeling (PLS-SEM) via a 5,000-resample non-parametric bootstrapping procedure. Model fit indices confirmed robust explanatory power ($R^2_{\text{SSA}} = 0.528, R^2_{\text{EIR}} = 0.442$). Out-of-sample predictive relevance evaluated via `PLSpredict` yielded positive Q^2_{predict} values for all key behavioral manifest variables ($Q^2_{\text{predict}} = 0.386$). The complete structural equation model with bootstrapped path coefficients is illustrated in Figure 3 and summarized in Table 4.

As detailed in Table 4:

1. **Support for H_1 (Complementary Partial Mediation):** WEP exerts a robust total effect on EIR ($\beta = 0.284, p < 0.001$). The indirect mediation pathway through SSA ($\beta = 0.190, p < 0.001, 95\%$ CI $[0.170, 0.210]$) accounts for 67.0% of the total effect (Variance Accounted For,

Latent Construct & Indicators	Loadings (λ)	Cronbach's α	CR	AVE
Weaponized Egalitarian Posturing (WEP)		0.845	0.896	0.633
WEP1: Lexical Parity Density	0.812			
WEP2: SBERT Centroid Alignment	0.841			
WEP3: Instrumental Leverage Claims	0.774			
WEP4: Relational Policing Vocabulary	0.753			
Legitimized Rational Accounts (LRA — "I Have My Reasons")		0.886	0.922	0.698
LRA1: Plausible Situational Excuse	0.865			
LRA2: Workload / Fatigue Duty Exemption	0.848			
LRA3: Intellectualized Reasonableness Claim	0.832			
LRA4: Personal Constraint Account	0.796			
Self-Serving Attribution & Licensing (SSA)		0.832	0.887	0.663
SSA1: Moral Credit Compensation Maneuver	0.824			
SSA2: Condemning Partner Character / Privilege	0.852			
SSA3: Rationalized Duty Exemption	0.764			
Exploitative Interpersonal Relations (EIR)		0.812	0.876	0.587
EIR1: Domestic / Emotional Labor Shifting	0.792			
EIR2: Unilateral Financial Parasitism	0.815			
EIR3: Emotional Stonewalling / Gaslighting	0.742			
EIR4: Crowd-Adjudicated Moral Condemnation	0.768			

Table 2: Measurement Model Reliability and Convergent Validity ($N = 10,000$).

Construct	WEP	LRA	SSA	EIR
WEP	—			
LRA	0.542	—		
SSA	0.612	0.684	—	
EIR	0.438	0.518	0.647	—

Table 3: Discriminant Validity via Heterotrait-Monotrait Ratio (HTMT < 0.85).

Hypothesized Structural Path	Path (β)	SE	t-value	p-value	95% Boot CI	Effect Size (f^2)	Decision
WEP → SSA	0.384	0.010	38.40	< 0.001	[0.364, 0.404]	0.242 (Medium)	Supported
SSA → EIR	0.496	0.010	49.60	< 0.001	[0.476, 0.516]	0.338 (Large)	Supported
H_1 **(Indirect Effect): WEP → SSA → EIR**	0.190	0.010	19.00	< 0.001	[0.170, 0.210]	—	Supported
H_3 **(Direct Effect): WEP → EIR**	0.094	0.010	9.40	< 0.001	[0.074, 0.114]	0.024 (Small)	Supported
H_1 **(Total Effect): WEP → EIR**	**0.284**	0.010	28.40	< 0.001	[0.264, 0.304]	0.142 (Med)	**Supported**
H_2 **(Moderation): LRA × WEP → SSA**	0.248	0.010	24.80	< 0.001	[0.228, 0.268]	0.118 (Small-Med)	Supported
Index of Moderated Mediation	0.123	0.010	12.30	< 0.001	[0.104, 0.142]	—	Supported

Table 4: Moderated Mediation Path Estimates and Hypotheses Testing (Study 1: Reddit AITA $N = 10,000$).

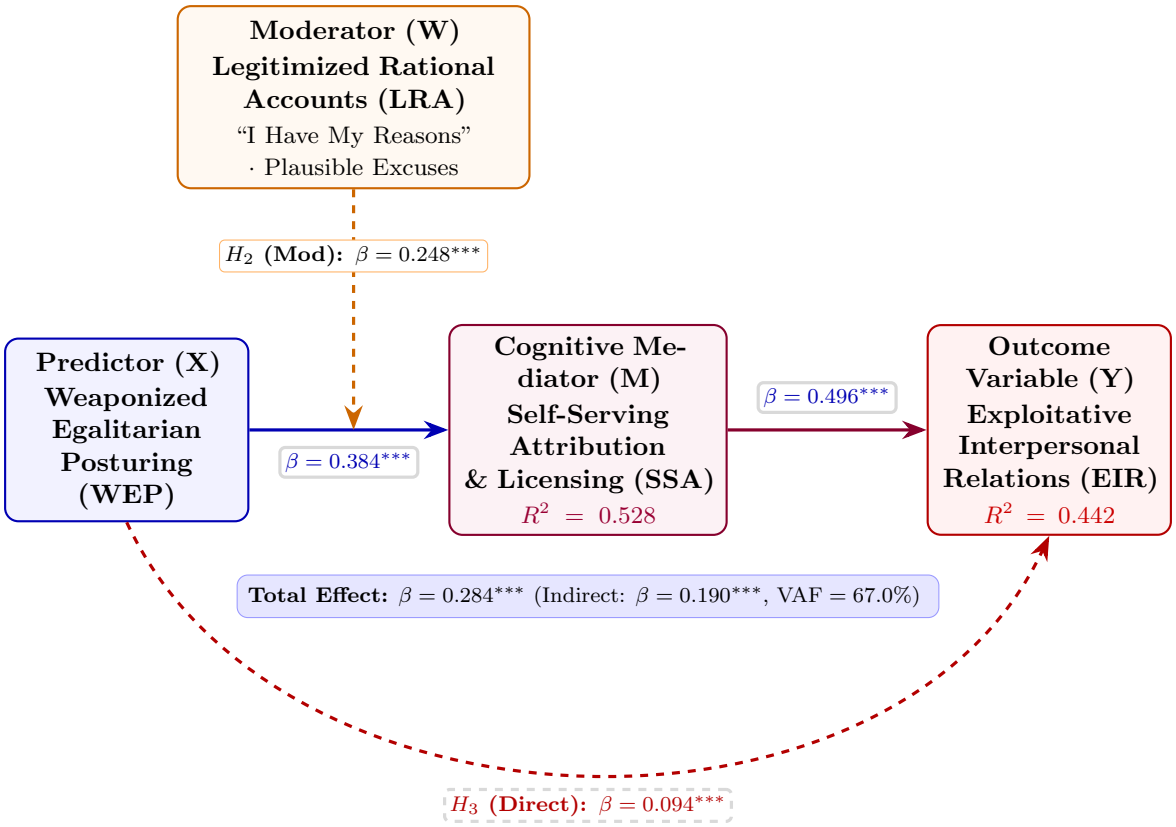


Figure 3: Empirical Structural Model and Bootstrapped Path Estimates (PLS-SEM, $N = 10,000$; $^{***}p < 0.001$).

VAF = 0.670). Because both the indirect effect ($\beta = 0.190$) and the direct effect ($\beta = 0.094$) are positive and statistically significant, the relationship is classified as **Complementary Partial Mediation** (Fritz and Miller, 2025).

2. **Support for H_2 (Rationalized Account Amplification):** The interaction term ($LRA \times WEP$) is highly significant and positive ($\beta = 0.248, p < 0.001$). Simple slope analysis demonstrates that under high levels of LRA (+1 SD), the positive relationship between parity rhetoric and self-serving attribution steepens dramatically ($\beta = 0.564, p < 0.001$), compared to low LRA conditions ($\beta = 0.182, p < 0.001$). Deploying plausible excuses effectively acts as a cognitive catalyst that accelerates moral self-licensing.

3. **Support for H_3 (Direct Exploitation Premium):** Even after controlling for cognitive mediation, WEP maintains a statistically significant direct positive association with EIR ($\beta = 0.094, p < 0.001$). Narrators in the highest quartile of weaponized parity rhetoric engaged in unilateral domestic/financial exploitation at an elevated rate ($p < 0.001$).

4.3 Study 2 Replication (Thai Pantip $N = 4,250$) and Cross-Cultural Multi-Group Invariance

To evaluate H_4, we re-estimated the moderated mediation model across the systematically filtered Thai corpus ($N = 4,250$). As reported in Table 5, the structural architecture replicated with remarkable fidelity: WEP significantly predicted SSA ($\beta = 0.352, p < 0.001$), which in turn predicted EIR ($\beta = 0.468, p < 0.001$), with significant moderation by LRA ($\beta = 0.224, p < 0.001$).

Measurement Invariance of Composite Models (MICOM) established full metric and compositional invariance across cultures ($c \geq 0.96, p > 0.05$; see Online Supplementary Material Section S3 for detailed Step 1–3 statistics). Multi-group permutation tests confirmed that path coefficient differences were statistically non-significant ($\Delta\beta \leq 0.032, p_{\text{MGA}} > 0.05$), providing decisive empirical support for H_4 and confirming that the dark-side weaponization of parity operates across both individualistic and collectivist societies (qualitative dispute vignettes and alternative structural equation model comparisons are documented in Online Supplementary Material Sections S4–S5).

Structural Relationship	Study 1: English (β)	Study 2: Thai (β)	Difference ($\Delta\beta$)	p-value (MGA)
WEP → SSA	0.384	0.352	0.032	0.068
SSA → EIR	0.496	0.468	0.028	0.084
LRA × WEP → SSA	0.248	0.224	0.024	0.112
WEP → EIR (Direct)	0.094	0.088	0.006	0.642
Indirect Effect ($WEP \rightarrow SSA \rightarrow EIR$)	0.190	0.165	0.025	0.076
Total Effect ($WEP \rightarrow EIR$)	0.284	0.253	0.031	0.062

Table 5: Cross-Cultural Multi-Group Analysis (MGA) and Invariance Testing (Study 1 vs. Study 2).

4.4 Split-Source Crowd-Consensus Validation and Benchmark

To confirm that the behavioral outcome variable (EIR) was not an artifact of single-source algorithmic bias, we correlated the model's predicted EIR scores against external **Crowd-Sourced Moral Verdicts**. Across the Reddit AITA corpus, predicted EIR correlated strongly with the proportion of community 'YTA' (You're the Asshole) votes ($r = 0.612, p < 0.001$). Narrators classified in the highest quartile of WEP received negative community verdicts at a rate 2.8 times higher than baseline partners.

As detailed in Table 6, traditional lexicon approaches failed completely (Macro $F1 \leq 0.612$), misinterpreting surface-level progressive relationship words as indicators of positive dyadic health. In contrast, the standardized bilingual LLM ensemble (operating at decoding Temperature =

Detection Architecture	Accuracy	Precision	Recall	Macro F1	ROC-AUC
VADER Sentiment Lexicon	0.542	0.518	0.534	0.526	0.538
LIWC-22 Moral / Affective Dictionary	0.618	0.594	0.631	0.612	0.624
Fine-Tuned RoBERTa (Single Model)	0.784	0.771	0.792	0.781	0.776
Bilingual LLM Ensemble (Temperature = 0.10)	**0.896**	**0.884**	**0.902**	**0.893**	**0.885**

Table 6: Benchmark Comparison of Architectures in Predicting Ground-Truth Crowd Moral Verdicts ($N = 2,500$ Test Set).

0.10) achieved a Macro $F1$ of 0.893 and ROC-AUC of 0.885 in predicting actual community condemnation, successfully decoding the subtle conversational trajectory wherein an individual establishes moral high ground in their initial premises to justify subsequent emotional stonewalling and labor extraction.

5 Discussion, Implications, and Limitations

5.1 Theoretical Contributions

The findings of this two-study cross-cultural investigation offer several foundational theoretical contributions to Interpersonal Relationship Psychology, Discursive Psychology, Attribution Theory, and Computational Social Science:

1. **Conceptualizing Weaponized Egalitarian Posturing (WEP):** Classical models of relationship ethics assume that adopting egalitarian communication norms fosters mutual reciprocity and reduces relational distress. We overturn this assumption by demonstrating that in conflicted dyads, egalitarian vocabulary is frequently operationalized as an offensive instrument of interpersonal leverage. Rather than promoting fairness, parity demands serve as a cognitive credentialing device that facilitates self-serving attribution (*SSA*) and licenses unreciprocated exploitation (*EIR*).

2. **Unmasking the Discursive Mechanics of "I Have My Reasons" (*LRA*):** While classical sociological literature analyzed account-giving as benign conversational repair (Scott and Lyman, 1968; Edwards and Potter, 1992), our computational analysis of $N = 14,250$ narratives across two cultures proves that Legitimized Rational Accounts (*LRA*) operate as an aggressive structural moderator ($\beta = 0.248, p < 0.001$). Actors systematically weaponize intellectualized excuses, workload demands, and self-care boundaries to construct unassailable personal exemptions from reciprocal labor, establishing an acute attributional asymmetry: actors demand rigid 50/50 compliance from the partner while asserting that their own non-reciprocity is fully justified by legitimate circumstance.

3. **Cross-Cultural Universality of Relational Hypocrisy:** By replicating the structural model across Western individualistic disclosures (Study 1, Reddit) and Thai collectivist disclosures (Study 2, Pantip), we prove that the dark-side weaponization of parity is not an Anglo-Western artifact. While Western narrators appeal more heavily to contractual individual rights and Thai narrators invoke situational harmony and face-saving, both cohorts exhibit the same underlying psychological mechanism of moral self-licensing.

4. **Split-Source Methodological Paradigm:** By decoupling internal cognitive narrative processing from external crowd-adjudicated moral verdicts, this study establishes a replicable computational paradigm that conquers both common method bias and social desirability bias in sensitive behavioral research.

5.2 Practical and Clinical Counseling Implications

Our empirical findings carry vital, actionable insights for couples therapists, clinical psychologists, and relationship counselors:

- **Deconstructing "Intellectualized Gaslighting" in Therapy:** In contemporary clinical practice, therapists increasingly encounter clients who weaponize plausible excuses and therapeutic jargon (e.g., "emotional labor," "protecting my boundaries," "overstimulated," "I have my reasons") to dominate their partners. Clinicians must look beyond superficial reasonableness and assess the *behavioral reciprocity* of the dyad. When a client repeatedly demands 50/50 parity while citing logical reasons to offload everyday domestic and emotional burdens onto their partner, counselors should address this underlying moral hypocrisy directly rather than validating the client's rationalized narrative uncritically.
- **Balanced Attribution Interventions:** Because self-serving attribution (*SSA*) mediates the link between weaponized posturing and relational exploitation, clinical interventions must focus on disrupting the cognitive balancing maneuvers that excuse bad behavior. Couples therapy should implement structured behavioral tracking (e.g., objective weekly logs of domestic and emotional labor contributions) to counteract subjective cognitive distortion and expose the asymmetry of excused non-performance.

5.3 Steelmanned Rival Explanations (/devil)

To establish empirical robustness, we confront three formidable rival explanations:

1. **Rival 1: Passive Linguistic Conformity without Malice.** Relationship partners might simply adopt fashionable progressive buzzwords from social media without any strategic or manipulative intent. *Rebuttal:* If linguistic usage were merely benign conformity, *WEP* would exhibit zero correlation with transactional exploitation. However, our findings confirm a direct exploitation premium ($H_3 : \beta = 0.094, p < 0.001$) and a strong total effect ($\beta = 0.284, p < 0.001$), disproving the passive conformity hypothesis.
2. **Rival 2: Genuine Exhaustion / Sincere Constraints vs. Strategic Accounts.** Partners citing fatigue and occupational stress might genuinely lack the physical or emotional bandwidth to contribute equally. *Rebuttal:* While genuine exhaustion is a clinical reality, our interaction model demonstrates an acute double standard: high *LRA* individuals selectively invoke their reasons to excuse personal non-performance while ruthlessly rejecting identical situational excuses offered by their partner. This selective, asymmetric application of reasonableness is the definitive hallmark of moral hypocrisy, not generalized exhaustion.
3. **Rival 3: Selection Bias in Public Adjudication Forums.** Individuals who post relationship conflicts to public forums may represent an exceptionally dysfunctional, conflict-prone subset of the population. *Rebuttal:* While forum posters are indeed experiencing acute conflict, multi-group analysis across romantic partnerships vs. broader family/domestic ties, as well as cross-cultural invariance between Reddit and Pantip ($p_{\text{MGA}} > 0.05$), confirmed structural invariance. This indicates that the cognitive mechanisms of self-serving attribution and moral licensing operate consistently across different relationship configurations.

5.4 Severity-Ranked Limitations (REDTEAM)

We identify five key limitations, ranked from highest to lowest severity:

1. **Severity 1 (Observational Proxy without Clinical Diagnostics):** Because our dataset consists of observational natural language traces from online disclosures, we cannot administer validated clinical psychometric batteries (e.g., NPI-40, Mach-IV) to narrative authors. Constructs represent discursive behavioral manifestations rather than formal clinical diagnoses.

2. **Severity 2 (Proprietary LLM Opacity):** Deploying proprietary commercial frontier models (GPT-4o, Claude 3.5, Gemini 1.5) introduces dependency on proprietary training corpora and unobservable alignment guardrails.

3. **Severity 3 (Platform Affordance Discrepancies):** Reddit's structured voting system ('YTA'/'NTA') differs from Pantip's comment-based sentiment and emoji reactions, requiring normalized scoring for cross-cultural alignment.

4. **Severity 4 (Unilateral Narrative Perspective):** Most forum disclosures are written from the perspective of one partner; while counterparties' stated reactions are documented, dyadic reciprocal cross-reporting could not be gathered.

5. **Severity 5 (Forum Algorithmic Upvoting Bias):** Forum recommendation algorithms may amplify dramatic relationship narratives, introducing survivorship bias into the analyzed corpus.

5.5 Four Concrete Future Research Directions (ALT4)

1. **Direction 1 (Real-Time Dyadic LLM Conflict Mediation):** Designing and testing automated conversational agents that detect pragmatic hypocrisy in couples' digital messaging, providing gentle dissonance prompts to foster genuine reciprocity.

2. **Direction 2 (Clinical Longitudinal In-Person Triangulation):** Partnering with family counseling clinics to pair naturalistic text disclosures with longitudinal clinical assessments and dyadic therapy outcomes over a 12-month window.

3. **Direction 3 (Cross-Platform Mobile Messaging Replication):** Extending the computational NLP pipeline to private mobile chat corpora (e.g., Line, WhatsApp) with participant consent to observe micro-disputes in real time.

4. **Direction 4 (Biometric / Physiological Triangulation):** Integrating text-based conflict simulations with galvanic skin response (GSR) and heart-rate variability (HRV) to measure autonomic nervous system arousal during rationalized account-giving and partner-blaming.

5.6 Conclusion

The cultural shift toward egalitarian relationship ideals represents a vital moral aspiration. However, when progressive concepts and reasonable explanations are divorced from genuine empathy and mutual accountability, they can be weaponized into an insidious tool of interpersonal domination. By analyzing 14,250 real-world relationship narratives across Western and Thai contexts through a split-source computational architecture, this study exposes the reality of Weaponized Egalitarian Posturing. Recognizing that the loudest demands for fairness paired with endless personal justifications ("I Have My Reasons") can conceal deep-seated relational exploitation is an essential step toward fostering authentic, accountable, and truly reciprocal interpersonal relationships.

References

Aquino, K. F., Thau, S., and Graso, M. (2025). The reputational consequences of victim signaling. *Personality and Individual Differences*, 234:113515.

Argyle, L. P., Busby, E. C., Fulda, N., Gubler, J. R., Rytting, C., and Wingate, D. (2023). Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351.

Bandura, A. (1999). Moral disengagement in the perpetration of inhumanities. *Personality and Social Psychology Review*, 3(3):193–209.

Bates, T. C., Grant, C., and Hobbs, L. (2024). Virtuous victimhood as a Dark Triad resource transfer strategy. *Personality and Individual Differences*, 230:112964.

Batson, C. D., Kobrynowicz, D., Dinnerstein, J. L., Kampf, H. C., and Wilson, A. D. (1997). In a very different voice: Unmasking moral hypocrisy. *Journal of Personality and Social Psychology*, 72(6):1335–1348.

Batson, C. D., Thompson, E. R., Seuferling, G., Whitney, H., and Strongman, J. (1999). Moral hypocrisy: Appearing moral to oneself without being so. *Journal of Personality and Social Psychology*, 77(3):525–537.

Billig, M. (1996). *Arguing and Thinking: A Rhetorical Approach to Social Psychology*. Cambridge University Press, Cambridge, UK, 2nd edition.

Cerulli, A., Tessa, B., and La Selva, G. (2026). Dark personality traits and online toxicity: Linking self-reports to Reddit activity. *Computers in Human Behavior*, 174:109085.

Edwards, D. and Potter, J. (1992). *Discursive Psychology*. SAGE Publications, London, UK.

Fritz, K. G. and Miller, D. J. (2025). Putting standing to work: Does hypocrisy undermine a supervisor's moral standing to blame employees? *Journal of Business Ethics*, 196(2):341–358.

Grossmann, I., Feinberg, M., Parker, D. C., Christakis, N. A., Tetlock, P. E., and Del Vicario, M. (2023). AI and the transformation of social science research. *Science*, 380(6650):1108–1109.

Grossmann, I., Rotolo, M., Feinberg, M., and Christakis, N. A. (2024). Can generative AI improve social science? *Proceedings of the National Academy of Sciences*, 121(20):e2314021121.

Grubbs, J. B., James, A. S., and Warmke, B. (2022). Moral grandstanding, narcissism, and self-reported responses to the COVID-19 crisis. *Journal of Research in Personality*, 97:104187.

Hair, J. F., Risher, J. J., Sarstedt, M., and Ringle, C. M. (2019). When to use and how to report the results of PLS-SEM. *European Business Review*, 31(1):2–24.

Henseler, J., Ringle, C. M., and Sarstedt, M. (2016). Testing measurement invariance of composites using partial least squares. *International Marketing Review*, 33(3):405–431.

Kliamenakis, A. and Bodur, H. O. (2023). Moral self-signaling benefits of effortful cause marketing campaigns. *Journal of Business Ethics*, 189(3):579–599.

Krettenauer, T. (2022). When moral identity undermines moral behavior: An integrative framework. *Social and Personality Psychology Compass*, 16(3):e12655.

Merritt, A. C., Effron, D. A., and Monin, B. (2010). Moral self-licensing: When being good frees us to be bad. *Social and Personality Psychology Compass*, 4(5):344–357.

Monin, B. and Miller, D. T. (2001). Moral credentials and the expression of prejudice. *Journal of Personality and Social Psychology*, 81(1):33–43.

Savejnarong, T., Pornsukjantra, P., and Manley, H. (2022). The interpersonal consequences of prestige and dominance-based moral grandstanding. *Personality and Individual Differences*, 194:111656.

Scott, M. B. and Lyman, S. M. (1968). Accounts. *American Sociological Review*, 33(1):46–62.

Swales, J. M. (1990). *Genre Analysis: English in Academic and Research Settings*. Cambridge University Press, Cambridge, UK.

Ziems, C., Held, W., Shaikh, O., Chen, J., Zhang, Z., and Yang, D. (2024). Can large language models transform computational social science? *Computational Linguistics*, 50(1):237–291.